Learning Travel-Time Embeddings for Real-Time Isochrone Queries

Mathias Vandaele

Independent Research
mathias.vandaele@pm.me

April 2026

Abstract

Computing travel times between addresses traditionally requires graph routing engines such as OSRM or Valhalla, whose per-query latency makes interactive applications — real-time isochrones, dense clustering, spatial joins over millions of points — prohibitive. We present a learned continuous representation of ∼33 million French addresses in which the Euclidean distance between two embedding vectors directly approximates their driving time. A single-tower network maps (lat, lon) through a Space2Vec multi-scale Fourier encoding, a small residual backbone with GELU pre-activations, and a linear projection into $\mathbb{R}^{64}$; we show that this linear head induces a learned Mahalanobis metric in backbone feature space. Trained against ≈432 million OSRM-supervised pairs with a pure ℓ_1 loss, the model reaches a mean absolute error of **1.40 minutes** (Spearman $\rho = 0.9989$), with 96.8% of predictions within five minutes of ground truth. A custom vector database, RingDB, answers *ring* queries $\{v : d_{\min} < \|q - v\| < d_{\max}\}$ over the full corpus via a SIMD-vectorized brute-force scan that exploits pre-cached squared norms, saturates the machine's memory bandwidth on commodity hardware, and outperforms a pivot-based pruning index we also tried. We discuss why, the surprising empirical finding that the embedding space spontaneously recovers France's road topology without ever seeing map data, and the negative result that made us abandon hyperbolic geometry for this task.

1 Introduction

Interactive mapping applications ask a simple question at enormous scale: "which places are reachable from here in less than T minutes?" Answering this well requires driving times, not straight-line distances, because real road networks are shaped by highways, one-way streets, rivers, mountains, and urban density. The canonical answer is to run a shortest-path query on a routing graph [5, 11]. For a single query this is fast; for *millions* of queries — the workload behind live isochrones on a slider, neighbourhood clustering, or dense spatial joins — it is not, and the engineering around precomputation, contraction hierarchies, and tile caches dominates production systems.

This paper explores a different trade-off. We ask whether driving time can be approximated well enough by a *continuous* function of a pair of coordinates that the routing graph can be replaced, at inference time, by a small neural network plus a vector database. Concretely, we want a mapping $f : \mathbb{R}^2 \to \mathbb{R}^d$ such that

$$\left\| f(x_1) - f(x_2) \right\|_2 \;\approx\; T(x_1, x_2), \tag{1}$$

where T denotes driving time. If such an f exists with sufficient accuracy, every question that used to require a routing engine collapses into a metric-space operation: "reachable in T minutes"

becomes a ball query, "closest reachable addresses" becomes k-nearest neighbours, "the boundary of the reachable region" becomes a ring query.

Our contributions are as follows.

1. **A compact model for a country-scale metric.** A ≈10 M-parameter single-tower network encodes ∼33 M French addresses into $\mathbb{R}^{64}$ such that Euclidean distance matches OSRM travel time to 1.40 minutes MAE on held-out pairs.

2. **A geometric interpretation of the projection head.** The final linear layer $W \in \mathbb{R}^{64\times 1024}$ is not an arbitrary bottleneck: we prove that it induces a learned Mahalanobis metric $M = W^\top W$ of rank at most 64 in the backbone's 1024-dimensional feature space. This explains why a purely linear head is sufficient and why further non-linearity there would be redundant with the backbone.

3. **RingDB: brute-force ring queries done right.** Instead of returning a full ball or a k-NN list, RingDB returns exactly the points in the annular shell $\{v : d_{\min} < \|q - v\| < d_{\max}\}$, which is the natural primitive for isochrone frontiers. Rather than pruning with pivots (which we tried and abandoned), RingDB pre-caches every vector's squared norm and reduces the entire query to one GEMV plus an elementwise test via the identity $\|a - b\|^2 = \|a\|^2 + \|b\|^2 - 2\langle a, b\rangle$. On commodity laptop hardware this approaches the machine's memory bandwidth ceiling.

4. **A negative result on hyperbolic geometry.** An earlier version of this work used Poincaré-ball embeddings, motivated by the intuition that road networks are hierarchical. We report in detail why this failed — learnable-curvature experiments let the model itself vote against the ball — and why Euclidean space with a learned Mahalanobis metric is a better fit for a road *mesh*.

5. **Emergent topology.** A UMAP projection of the learned embeddings reveals that the model has spontaneously reconstructed the French road network — highway filaments, geographically isolated clusters, physical barriers such as rivers and rail lines — despite never having seen a map. Travel-time supervision alone is sufficient to recover topology.

2 Related Work

Graph-based routing. Production routing engines such as OSRM [5] and Valhalla [11] implement contraction hierarchies and multi-level Dijkstra variants over OpenStreetMap. They answer single queries in milliseconds but scale badly to workloads that demand many simultaneous reachability computations; approximate alternatives such as time-dependent contraction hierarchies still need the graph in memory.

Learned ETA. Recent work at production scale has used graph neural networks and sequence models to predict estimated times of arrival conditional on a given route [2]. These models assume the route already exists and focus on traffic-aware adjustment. Our setting is different: we want the pairwise time between *arbitrary* origins and destinations, without ever materializing a route.

Metric embeddings. Learning embeddings such that a simple distance approximates a costly target function is a classical idea in information retrieval and metric learning [4, 9]. The closest cousin to our approach is work that learns Euclidean embeddings of graph nodes so that shortest-path distances can be approximated by a simple vector operation [8], but that line of work operates on the nodes of a fixed, given graph — typically a social network — and does not

ingest raw geographic coordinates. We want the opposite: a model that takes (lat, lon) in and produces an embedding out, with no graph ever materialised at inference time.

Coordinate encodings. Fourier feature encodings from NeRF [10] and Space2Vec [6] have become the default way to feed low-dimensional continuous positions into MLPs. We use the Space2Vec variant — log-spaced wavelengths covering five orders of geographic scale, combined with random unit directions — as the input block of our network.

Hyperbolic representations. Poincaré-ball embeddings excel at tree-like hierarchies [3, 7]. Early in this project we conjectured that road networks, with Paris as a root and departments as branches, would benefit from such geometry. Our Section 8 reports why this intuition is wrong for driving times: real road networks are *meshes*, not trees, and hyperbolic space collapses everything to the ball's boundary.

3 Method

3.1 Overall architecture

The network is a single tower that maps a pair of geographic coordinates to a vector in $\mathbb{R}^{64}$:

$$(\text{lat}, \text{lon}) \longrightarrow \text{xyz} \longrightarrow \underbrace{\phi(\cdot)}_{\text{Space2Vec}} \longrightarrow \underbrace{g(\cdot)}_{\text{residual backbone}} \longrightarrow \underbrace{W \cdot (\cdot) + b_W}_{\text{linear head}} \longrightarrow z \in \mathbb{R}^{64}.$$

Since it is a single tower, the induced metric is symmetric by construction: $\|f(A) - f(B)\| = \|f(B) - f(A)\|$. An earlier dual-tower variant was explored to capture asymmetries caused by one-way streets and elevation, but on the static, daytime-averaged OSRM labels used here the symmetric tower matches it in accuracy and is simpler to deploy. Total trainable parameters: 10,057,281 (backbone 9,974,784; projection head 65,600; mapping MLP 16,897; Space2Vec buffer frozen). Figure 1 shows the full pipeline.

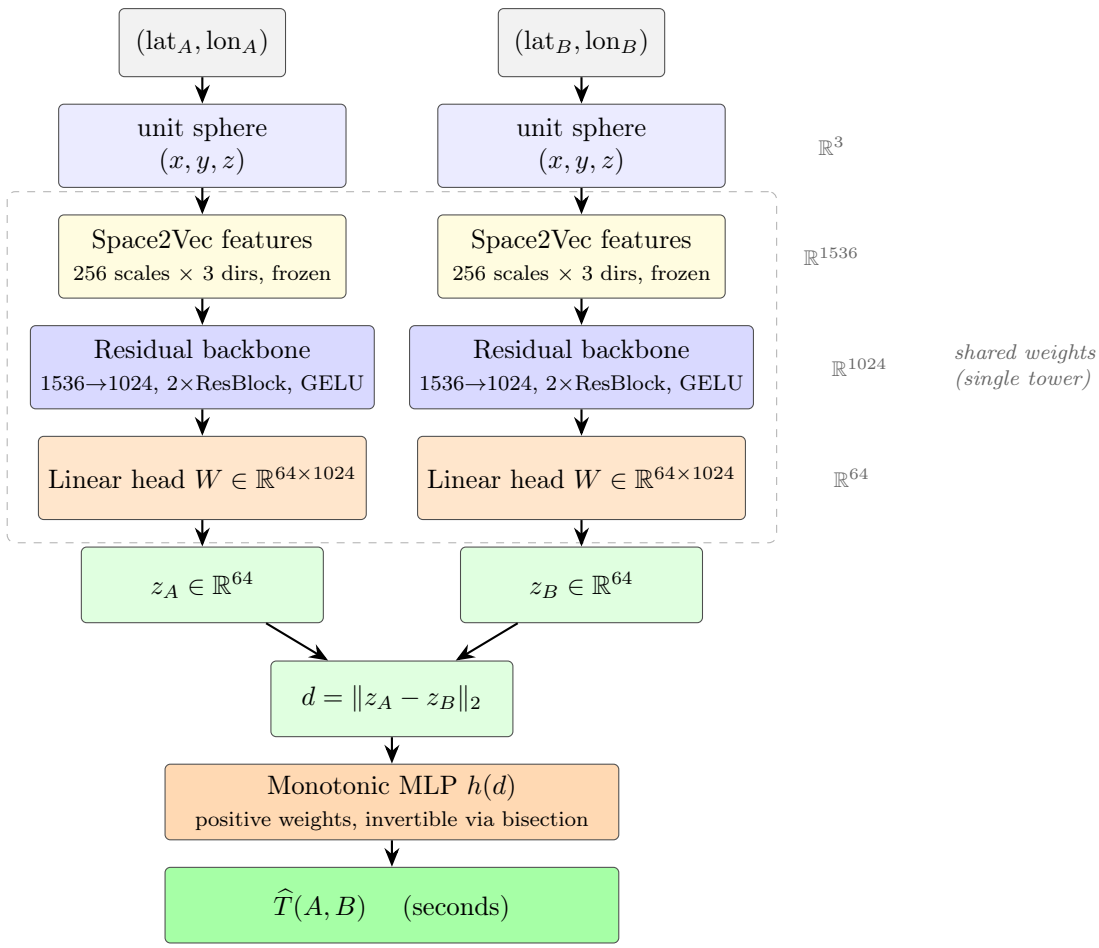

Figure 1: The pipeline. Two addresses A and B are independently encoded by a *single* shared tower — Space2Vec multi-scale Fourier features, a residual backbone, and a linear projection head into $\mathbb{R}^{64}$ — and their embeddings are compared by Euclidean distance. A monotonic MLP then maps this distance to a predicted travel time in seconds. Because the mapping is monotone, it is invertible by bisection, which turns isochrone queries "reachable in T minutes" into ring queries in embedding space. The linear head W induces a learned Mahalanobis metric $M = W^\top W$ of rank at most 64 in the backbone's 1024-dimensional feature space (Proposition 1).

3.2 Input: unit-sphere coordinates (no centering)

Latitude and longitude are mapped to unit-sphere coordinates $(x, y, z) = (\cos\phi\cos\lambda, \cos\phi\sin\lambda, \sin\phi)$ to remove the longitudinal compression of planar projections. No recentering or rescaling is applied. Naively, one might worry that metropolitan France occupies only ∼0.04% of the unit sphere, giving Fourier features too narrow an input range; Space2Vec's multi-scale wavelengths (below) handle this directly by placing basis functions at every geographically relevant scale, so the raw unit-sphere coordinates are sufficient.

3.3 Space2Vec Fourier features

We use a multi-scale Fourier encoding in the style of Space2Vec [6]. We choose $S = 256$ wavelengths log-spaced from $\lambda_{\min} = 10$ m to $\lambda_{\max} = 1000$ km,

$$\lambda_s \;=\; \lambda_{\min} \cdot \big(\lambda_{\max}/\lambda_{\min}\big)^{s/(S-1)}, \qquad s = 0, \ldots, S-1,$$

covering five orders of geographic scale, from city-block granularity to country-level structure. For each scale s we draw three random unit directions $\hat{b}_{s,1}, \hat{b}_{s,2}, \hat{b}_{s,3} \in \mathbb{R}^3$ and form frequency

vectors $b_{s,j} = \hat{b}_{s,j} \cdot (R_\oplus/\lambda_s)$, where $R_\oplus$ is the Earth radius. With this scaling, a geodesic of length λ_s advances the argument of $b_{s,j}^\top x$ by 2π, which is exactly what "a wavelength" means at the corresponding spatial scale. Stacking all $3S = 768$ frequency vectors as columns of a matrix $B \in \mathbb{R}^{3\times 768}$, the features are

$$\phi(x) \;=\; \big[\sin(2\pi\, B^\top x),\; \cos(2\pi\, B^\top x)\big] \in \mathbb{R}^{1536}.$$

The matrix B is stored as a non-trainable buffer. This resolves the dynamic-range problem without any input centering or rescaling: each wavelength contributes features tuned to its own geographic scale, and the model sees block-level, city-level, regional, and country-level structure from the first layer onward.

3.4 Residual backbone

The backbone g is a small residual MLP. An input projection takes $\mathbb{R}^{1536}$ to $\mathbb{R}^{1024}$, followed by two pre-activation residual blocks and a final LayerNorm. Each block implements

$$x \;\mapsto\; x \;+\; L_2\big(\,\mathrm{GELU}\big(L_1(\mathrm{GELU}(\mathrm{LN}(x)))\big)\,\big),$$

with a $2\times$ expansion ($1024 \to 2048 \to 1024$). The backbone has 9,974,784 parameters and is by far the largest component. The width of 1024 is deliberately chosen to be much larger than the output dimension: the backbone is an over-parameterized feature space that the linear head will then *select from*.

3.5 The linear projection head as a learned Mahalanobis metric

The final layer is a single projection $W \in \mathbb{R}^{64\times 1024}$ with bias b_W,

$$f(x) \;=\; W\, g(\phi(x)) \;+\; b_W.$$

The squared Euclidean distance between two embeddings is bias-free (the bias cancels in the difference) and factors cleanly:

$$\begin{aligned}\|f(A) - f(B)\|_2^2 &= \big\|W\,(g(A) - g(B))\big\|_2^2 \\ &= \big(g(A) - g(B)\big)^\top \underbrace{W^\top W}_{=:\,M} \big(g(A) - g(B)\big). \qquad (2)\end{aligned}$$

Proposition 1. *The matrix $M = W^\top W \in \mathbb{R}^{1024\times 1024}$ is symmetric positive semi-definite with* $\mathrm{rank}(M) \le 64$. *The squared embedding distance (2) is therefore a* Mahalanobis pseudo-distance *in the 1024-dimensional backbone feature space, with precision matrix M.*

Proof. Symmetry: $M^\top = (W^\top W)^\top = W^\top W$. For any v, $v^\top M v = \|Wv\|_2^2 \ge 0$, so M is PSD. Finally, $\mathrm{rank}(M) = \mathrm{rank}(W^\top W) = \mathrm{rank}(W) \le 64$. The bilinear form (2) therefore has the algebraic signature of a Mahalanobis metric, degenerate on $\ker(W)$. □ □

This observation has three practical consequences.

(i) Why a linear head is enough. A metric is a bilinear form; linear projection composed with ℓ_2 is exactly that form. Any further non-linearity in the head would *not* give us a different metric structure, only a different non-linear embedding — which the backbone can already express. The head is not an arbitrary bottleneck; it is the minimal functional form capable of producing a metric.

(ii) What the model actually learns. A rank-64 projection of a 1024-dimensional feature space picks out the 64 most travel-time-discriminative directions in backbone space. Directions

orthogonal to range($W^\top$) are explicitly ignored: the model says "these variations in the backbone's representation do not matter for travel time." The backbone is allowed to be rich and redundant; the head is what disciplines it.

(iii) Low effective dimensionality is predicted, not a bug. Empirically, the 64-dimensional embedding space has an intrinsic dimension of ≈9.9 (Section 5). Rather than treating this as "wasted capacity," Proposition 1 explains it: road-network travel times lie on a low-dimensional manifold, and the Mahalanobis metric discovers that dimensionality on its own. We return to this in Section 6, where it matters for the vector index.

3.6 From distance to seconds: a monotonic mapping MLP

The relationship between Euclidean distance in embedding space and travel time in seconds is monotonic but not linear: short urban trips and long inter-regional trips live in different regimes, and no two-parameter formula (linear, log, power, tanh) fits both well. We therefore use a small *monotonic* MLP $h : \mathbb{R}_{\geq 0} \to \mathbb{R}_{\geq 0}$ with three layers of width 128 and positive-weight reparameterization:

$$h(d) \;=\; \mathrm{softplus}\Big(W_3^+ \,\mathrm{ReLU}\big(W_2^+ \,\mathrm{ReLU}(W_1^+ d + b_1) + b_2\big) + b_3\Big),$$

where each weight matrix $W_i^+ = \mathrm{softplus}(\widetilde{W_i})$ is elementwise non-negative. Positive-weight monotone networks have a long history going back at least to Sill's min-max networks and the analysis of Daniels and Velikova [1]; more recent work by Wehenkel and Louppe [13] establishes that the positive-weight construction remains a universal approximator of monotone functions when composed with suitable activations. The composition of non-negative linear maps, ReLUs, and a final softplus is non-decreasing, which matters for two reasons:

- The prediction $\hat{T} = h(\|f(A) - f(B)\|)$ is automatically monotone in embedding distance, ruling out pathological solutions where nearby embeddings are predicted to be far in time.
- h is easily inverted by bisection in ≈20 iterations (cost < 0.01 ms). Isochrone queries — "addresses within T minutes" — are converted into ring queries by computing $h^{-1}(T \pm \varepsilon)$ once per request.

The mapping MLP has 16,897 parameters and is optimized jointly with everything else, with a 2× learning-rate multiplier.

4 Training

4.1 Dataset: 432 M OSRM-supervised pairs

Training pairs are generated from the French *Base Adresses Nationale* (BAN), which lists ≈33 M addressable points across metropolitan France. For each pair of addresses (A, B) we query OSRM with the default car profile for the driving time $T(A, B)$ in seconds. We generated ≈469.5 M raw pairs on a 112-thread AMD EPYC instance over seven hours on spot pricing (total cloud cost: $9); after removing overseas territories (poor OSM coverage) and near-zero trivial pairs, 431,789,950 records remain. The empirical time distribution is log-normal, with median 65 min and a tail reaching ∼14 h for cross-country trips.

We hold out 1 M pairs from the end of the shuffled Parquet file for validation and use the remaining ≈430 M for streaming training.

4.2 Loss

We minimise a pure ℓ_1 loss on predicted seconds:

$$\mathcal{L} \;=\; \frac{1}{N}\sum_{i=1}^{N}\big|\,\hat{T}_i - T_i\,\big|, \qquad \hat{T}_i = h\big(\|f(A_i) - f(B_i)\|_2\big).$$

Earlier variants used log-space MSE (to equalise relative error across scales) and smooth-ℓ_1; pure ℓ_1 gave uniformly better per-bucket MAE in our setting, likely because it neither under-penalizes long trips (as log-space MSE does) nor introduces a quadratic region near zero that biases the optimizer toward many small corrections instead of the large corrections the model actually needs in early training.

4.3 Optimisation

We train with AdamW, initial learning rate 10^{-3} (with a 2× multiplier for the mapping MLP), weight decay 10^{-5}, cosine annealing schedule with $\eta_{\min} = 10^{-6}$, batch size 8192, mixed precision (AMP) on a single NVIDIA GPU. Gradient clipping is set to $\|g\| \leq 5$. Ten epochs over the full 432 M dataset take ≈6 hours on an L4-class GPU; each epoch is ≈52,586 batches, with one validation pass per epoch.

5 Results

5.1 Accuracy

Table 1 reports headline metrics on 1 M held-out validation pairs after ten epochs. Mean absolute error is 1.40 minutes overall, with Spearman correlation $\rho = 0.9989$, meaning the induced ordering of reachable destinations is essentially perfect.

Table 1: Headline validation metrics after 10 epochs on 1 M held-out pairs.

Metric	Value
MAE (minutes)	**1.40**
MAE (seconds)	84
Pearson r	0.9659
Spearman ρ	**0.9989**
Predictions within 1 minute	52.8%
Predictions within 2 minutes	77.6%
Predictions within 5 minutes	96.8%
Predictions within 10 minutes	99.7%
MAPE (trips > 5 min)	3.9%
Median relative error	1.5%

Table 2 breaks this down by trip duration. Error scales sub-linearly: a 0–10 min trip has median error ≈0.6 min (≈36 s), while a 4 h+ trip has median error 1.5 min despite being on the order of 240 minutes long. This is the behaviour a log-normal error model would predict, but achieved here with a pure ℓ_1 loss.

5.2 Embedding geometry and intrinsic dimension

After training, embeddings have norm 31.2 ± 1.2 and mean pairwise L2 distance ≈7.05. A PCA on the validation embeddings shows that 24 components explain 90% of the variance and 57 explain 99%; the participation-ratio effective dimension is ≈9.9. Rather than a symptom of

Table 2: Per-bucket validation MAE at the best epoch.

Bucket	MAE (min)	Median (min)	P90 (min)	# pairs
0–10 min	0.8	0.6	1.7	35,317
10–30 min	1.4	1.0	3.0	26,809
30–60 min	1.4	1.0	3.0	31,894
1–2 h	1.4	1.0	3.0	58,124
2–4 h	1.8	1.1	4.1	34,925
4 h+	2.3	1.5	5.4	12,931

under-utilisation, this is *consistent* with Proposition 1: the Mahalanobis metric $M = W^\top W$ has rank at most 64, and the observed effective dimension reflects the intrinsic dimensionality of the travel-time manifold itself, which is much smaller than the nominal 64 available.

5.3 Emergent road topology

Projecting the 64-dimensional embeddings to 2D via UMAP reveals, without any geographic supervision, the structure of the French road network (Figure 2): major roads appear as thin filaments, cities cluster by commune, and travel-time barriers are visible as gaps and discontinuities between clusters.

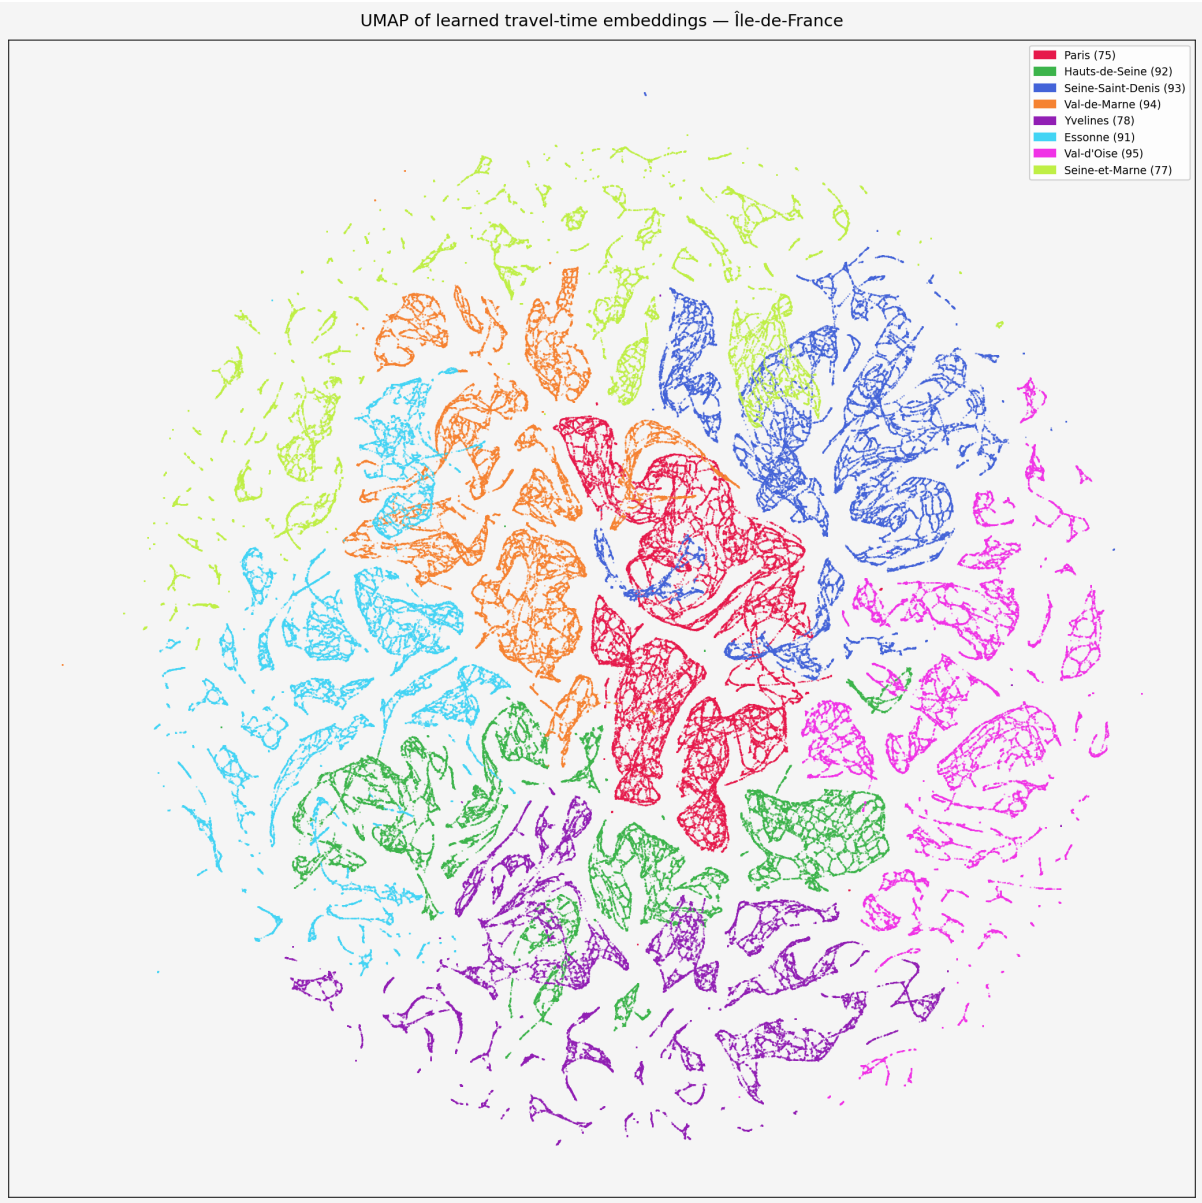

Figure 2: UMAP projection of learned 64-dimensional embeddings for addresses in Île-de-France, coloured by département. Paris (75, red) sits at the centre; the inner-ring départements — Hauts-de-Seine (92), Seine-Saint-Denis (93), Val-de-Marne (94) — form a tight cluster around it; the outer départements — Yvelines (78), Essonne (91), Val-d'Oise (95), Seine-et-Marne (77) — occupy the periphery. This concentric structure mirrors the actual geography of the Paris ring roads (Périphérique, A86, Francilienne) and was never provided to the model: it emerges purely from travel-time supervision. Within each département, filament structures correspond to individual road segments and branching points to junctions. Administrative boundaries are respected without ever appearing in the training signal.

6 RingDB: Brute-Force Ring Queries via Cached Norms

Accurate embeddings are only half of the story; production applications demand fast spatial queries over the full ∼33 M address corpus. We built RingDB [12], a specialised vector database whose primary query is not k-nearest neighbours but the annular *ring*:

$$\mathrm{Ring}(q, d_{\min}, d_{\max}) \;=\; \{\, v \in V \;\mid\; d_{\min} \,<\, \|q - v\|_2 \,<\, d_{\max} \,\}. \tag{3}$$

RingDB returns exactly the frontier of an isochrone without ever materialising its interior.

6.1 Why rings, not balls

Isochrone rendering does not need every address inside the reachable region — only the boundary. A contour in the plane is determined by its frontier points; interior points are redundant. Choosing $d_{\min} = h^{-1}(T - \varepsilon)$ and $d_{\max} = h^{-1}(T + \varepsilon)$ returns only the ∼5000 addresses lying on the edge of the "T-minutes reachable" region. Crucially, this keeps the result-set size roughly *constant* as the underlying corpus grows, which is what makes constant-latency isochrones possible.

6.2 The expansion identity

The design of RingDB turns on a single elementary observation. For any two vectors $a, b \in \mathbb{R}^d$,

$$\|a - b\|_2^2 \;=\; (a - b) \cdot (a - b) \;=\; \|a\|_2^2 \;+\; \|b\|_2^2 \;-\; 2\,\langle a, b\rangle. \tag{4}$$

If $\|v\|_2^2$ is precomputed and cached for every database vector, and $\|q\|_2^2$ is computed once per query, then the only per-vector work needed to evaluate $\|q - v\|_2^2$ is one inner product $\langle q, v\rangle$. Collecting all N database vectors as the rows of a matrix $D \in \mathbb{R}^{N \times d}$, the vector of inner products across the entire corpus is a single matrix–vector product, $D\,q \in \mathbb{R}^N$, and the full set of pairwise squared distances is

$$\text{sq_dists} \;=\; \underbrace{\|q\|_2^2\,\mathbf{1}}_{\text{scalar broadcast}} \;+\; \underbrace{n_v}_{\text{cached vector in } \mathbb{R}^N} \;-\; 2\,(D\,q),$$

where $n_v \in \mathbb{R}^N$ is the cached squared-norm column. The ring test (3) then reduces to a single elementwise comparison against $[d_{\min}^2,\, d_{\max}^2]$.

This turns the entire query into one GEMV (general matrix–vector product) plus an elementwise pass over a vector of length N. Both operations are perfectly sequential in memory, SIMD-friendly, and memory-bandwidth bound — the best-case workload for any modern CPU or GPU. In practice the implementation fuses the two operations into a single streaming pass over the rows of D (Section 6.4) so that the intermediate Dq vector is never materialised; the fused version is what the latency numbers below refer to.

6.3 What we tried before this, and why it failed

Our first design was a pivot-based index using the triangle inequality in the style of classical metric trees. A small set of k-means centroids served as pivots; for each data point we stored its distance to every pivot, and at query time we used the triangle inequality to bound $\|q - v\|$ and prune candidates without ever loading their coordinates. On paper this looked like a clean win: real embeddings have low intrinsic dimension (≈9.9), which is exactly the regime where pivot pruning is supposed to shine.

In practice, *almost nothing got pruned.* The embedding distribution is geometrically tight — norms 31.2 ± 1.2, mean pairwise distance ≈7 — so triangle-inequality bounds rarely excluded candidates with high confidence, and the per-candidate overhead of bitmap lookups and bound evaluations was not cheap enough to amortise against the few candidates that *were* eliminated. The end-to-end query time was consistently worse than a straightforward brute-force scan.

The brute-force scan via (4) has two decisive advantages that pivot methods struggle to match:

1. **No per-candidate branching.** Every database point is touched in the same way, via fused multiply–adds in a tight dot-product loop. There is no cache-hostile random access,

no bitmap intersection, no bound evaluation. The inner loop is a dense reduction whose SIMD instruction mix has been a compiler target for decades and which a modern CPU can essentially run at memory-bandwidth speed.

2. **Pre-cached norms convert a quadratic problem into a linear one.** Without caching, $\|q - v\|^2$ would require reading both q and v, subtracting, and re-squaring per candidate. With cached $\|v\|^2$, only the dot product $\langle q, v\rangle$ is computed per candidate, and the rest is scalar arithmetic. The memory traffic becomes essentially $O(N \cdot d)$ (a single streaming read of D) plus $O(N)$ (a streaming read of cached norms), which is close to the minimum any algorithm could achieve for this problem.

6.4 Implementation

RingDB is implemented in Rust [12]. The on-disk layout is a handful of memory-mapped flat binary files: `meta.bin` records the pair (dims, N) as two little-endian `u64` values; `vectors.bin` holds the matrix D as row-major `f32` ($N \cdot d$ floats, ≈7.9 GB for $N = 33$ M at $d = 64$); `norms_sq.bin` holds the pre-computed squared norms $\|v_i\|_2^2$ as a contiguous `f32` column (≈130 MB); and `payloads.bin` plus `offsets.bin` store per-vector metadata (addresses, lat/lon, external IDs). All files are opened once via `mmap` and accessed as typed slices, so there is no deserialisation cost and no per-query allocation.

The core query primitive is a single parallel pass over the rows of D. For each row v_i the loop body computes one dot product $\langle q, v_i\rangle$, reconstructs $\|q - v_i\|_2^2$ from the cached norms via equation (4), and tests the result against $[d_{\min}^2, d_{\max}^2]$. A disk (ball) query is the same loop with only the upper bound checked; both primitives share the same scan code and differ only in a single constant-time comparison.

The dot product itself is a hand-written Rust routine with *four independent* floating-point accumulators fed by fused multiply–adds (`f32::mul_add`). Breaking the reduction into four parallel partial sums removes the loop-carried dependency that would otherwise bottleneck a naive dot product on FMA latency, which is essential for reaching peak floating-point throughput on modern pipelined FPUs. Combined with `chunks_exact(4)` over the input slices, this pattern is one that LLVM reliably auto-vectorises into `fmla.4s` on ARM NEON and `vfmadd` on x86 AVX; we verified the generated assembly with `cargo asm` to confirm that the intended SIMD instructions are actually emitted, and not replaced by a scalar fallback or broken up by the optimiser. No BLAS dependency is required: the compiler-emitted kernels match a hand-tuned GEMV closely enough that the additional integration cost would be wasted.

Parallelism across rows is handled by Rayon's `par_chunks_exact`: the scan is embarrassingly parallel along N, each worker thread owns a contiguous slab of the memory-mapped matrix, and the surviving candidates are gathered into the result via a standard parallel `filter_map`. The result type is a plain `#[repr(C)]` struct holding the vector index and squared distance — a *plain old datum* that can be memcopied and streamed over a wire without serialisation.

6.5 Measured performance

On a single Apple M3 Pro laptop, the exact brute-force scan over 30 M 64-dimensional vectors runs at ≈67 ms. This is within a small constant of the theoretical memory bandwidth ceiling of the machine: the scan is bandwidth-bound, not compute-bound, which means additional optimisation effort on the arithmetic side would be wasted. The full isochrone pipeline — embed the query address, invert h to obtain $(d_{\min}, d_{\max})$, run the ring scan, compute an alpha shape over the ∼5000 returned frontier points, serialise a GeoJSON polygon — completes well inside the 16 ms budget for 60 fps on modest server hardware, and in tens of milliseconds on the laptop. A GPU backend via CUDA is in progress and is projected to bring the scan to ∼8 ms on an RTX-class GPU through the same dot-product-plus-mask recipe.

7 Applications

The combination of an accurate metric and a ring index unlocks a small family of interactive applications that are awkward or impossible on top of a conventional routing engine.

A motivating example. Consider a real-estate search with compound reachability constraints: "show me apartments within 25 minutes of Gare de Lyon, within 45 minutes of Charles de Gaulle airport, and within one hour of Disneyland Paris." In the routing-engine framing this is a cross product of three shortest-path computations per candidate, over tens of millions of candidates — effectively impossible interactively, and cache-hostile at any scale. In our framing, each reachability constraint is a single disk query in embedding space (the same primitive RingDB uses for rings, with only the upper bound checked); the three result sets are intersected by ID; and the entire query, over the full 33 M-address corpus, completes in well under half a second on the laptop. The compound version of the question is not just *faster* than the routing-engine version — it is *practically new*, in the sense that no end user would ever have considered asking it of a routing API. That is what we mean when we say the metric embedding unlocks applications rather than merely accelerating them.

Real-time isochrones. A slider that controls "T minutes from here" becomes trivial: each drag event issues a ring query for the new T, runs a morphing concave hull on the $\sim$5000 returned frontier points, and streams the resulting polygon to a web client over WebSocket. Because the result-set size does not grow with T, latency is constant in the slider position.

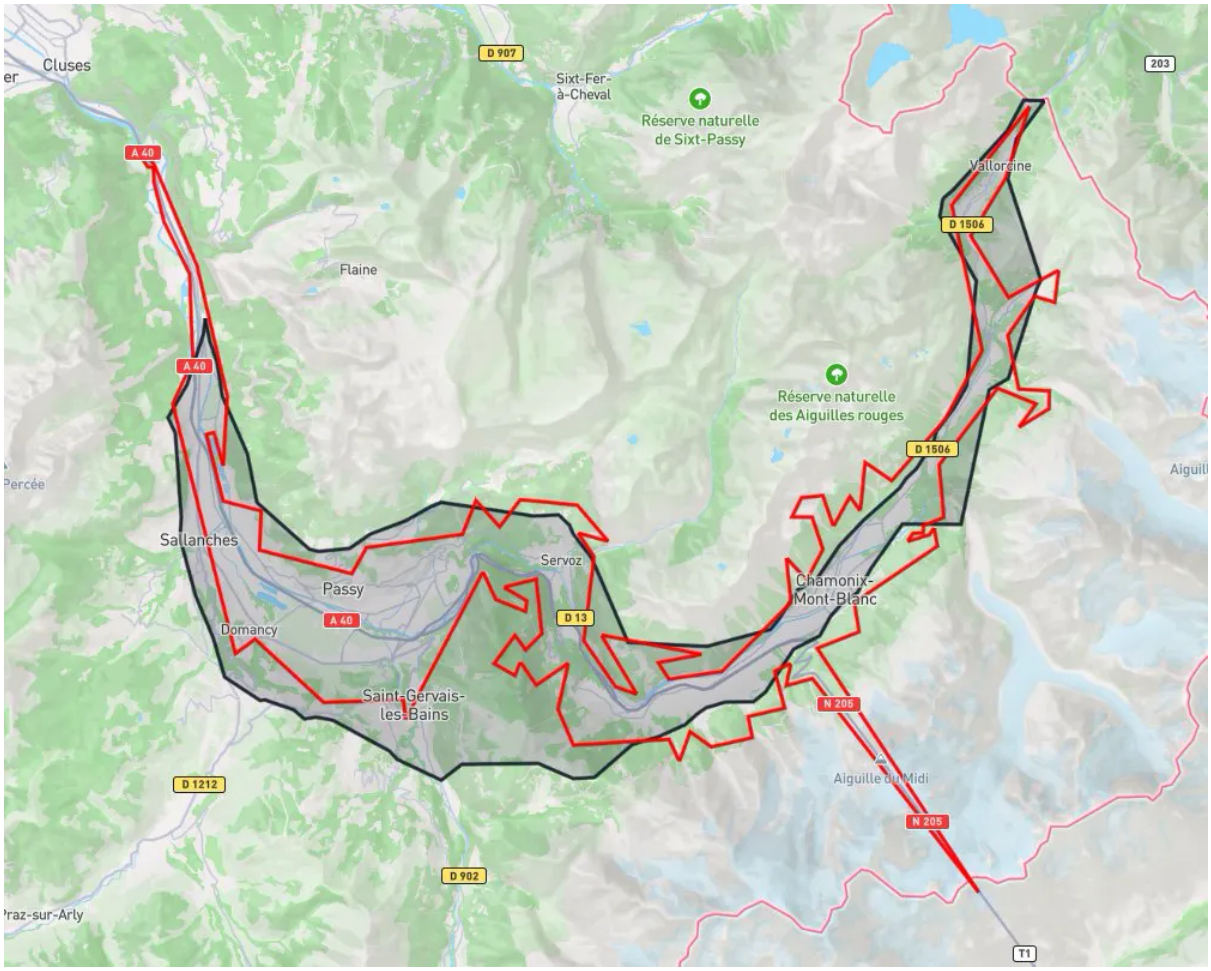

Figure 3: 45-minute driving isochrone centred on the Chamonix valley. **Black**: our system (ring query + alpha-shape contour). **Red**: OSRM ground truth. The valley crescent is correctly recovered despite the model never having seen map data during training. Differences arise from two expected sources: highway segments (A40, N205) are absent from the BAN address corpus and therefore produce no frontier points, eliminating the ground-truth spikes; and our contour smooths over micro-topology artifacts — switchbacks, dead-end mountain roads — that cause the OSRM polygon to be non-convex in ways that would be misleading on a rendered map layer. The Chamonix valley is a deliberate stress test: vertical Alpine walls and a single access corridor make this one of the most constrained isochrone geometries in metropolitan France.

Spatial clustering on travel time. Clustering addresses by driving-time similarity rather than straight-line distance is a natural request for logistics and location intelligence, and is intractable with a routing backend. In our embedding space it is a standard k-means or HDBSCAN over the 64-dimensional vectors.

Reachability joins for real estate. Listings can be filtered by compound constraints such as "within 25 minutes of a school, 40 minutes of an office, and 1 hour of an airport" by intersecting three ball queries in embedding space — a primitive RingDB supports directly by replacing the lower bound with zero.

Travel-time Voronoi diagrams. Given a set of points of interest, partitioning the map by closest-in-time POI becomes a single vector-index pass.

8 The Negative Result: Why Not Hyperbolic Space?

An earlier generation of this work embedded addresses into the Poincaré ball rather than Euclidean space. The intuition was seductive: road networks look hierarchical (Paris as a root, departments as branches), and hyperbolic space is famously efficient for trees [7]. We abandoned the approach after extensive experimentation because it failed in an instructive way.

With a fixed ball of curvature $c = 1$, embeddings systematically collapsed to the ball's boundary: $\|f(x)\| \to 1$ with standard deviation $\to 0$ across the dataset. Replacing the soft projection with a hard clamp at 0.9 simply produced collapse at 0.9. Adding a learnable curvature parameter let the model itself vote: c drifted from 0.1 to 6.57, i.e., the network chose to *shrink* the ball so that radius $R = 1/\sqrt{c}$ became tiny, which is mathematically equivalent to destroying the hyperbolic structure. On comparable 10 M-pair experiments, Euclidean embeddings reached ≈14.5 min MAE while the best Poincaré variant stalled at ≈14.9 min with degenerate boundary norms.

The diagnosis we settled on is that *France's road network is a mesh, not a tree.* Lille–Strasbourg, Lyon–Marseille, Bordeaux–Toulouse are all *direct* connections that do not route through any central hub. Mesh metric spaces need uniform capacity at all scales, which Euclidean (and a Mahalanobis metric over Euclidean) provides and which the hyperbolic ball does not: in d dimensions the hyperbolic volume element scales as $\sinh^{d-1}(r)$, so almost all representational capacity is concentrated in a vanishing shell near the boundary, and gradient descent pushes everything there. The learnable-curvature experiment is, in retrospect, the cleanest possible evidence: given the option, the model itself votes against the ball.

Figure 4 visualises this trajectory on the 10 M-pair experiment. Initialised at $c = 0.1$ (a Poincaré ball of radius $R = 1/\sqrt{c} \approx 3.16$), gradient descent steadily pushes c upward to 6.57 over twenty epochs, shrinking the ball's radius to $R \approx 0.39$. Throughout, the embedding-norm standard deviation stays pinned at zero: *every* embedding sits on the boundary regardless of where that boundary is. The model is not finding a better ball; it is trying to destroy the one it was given.

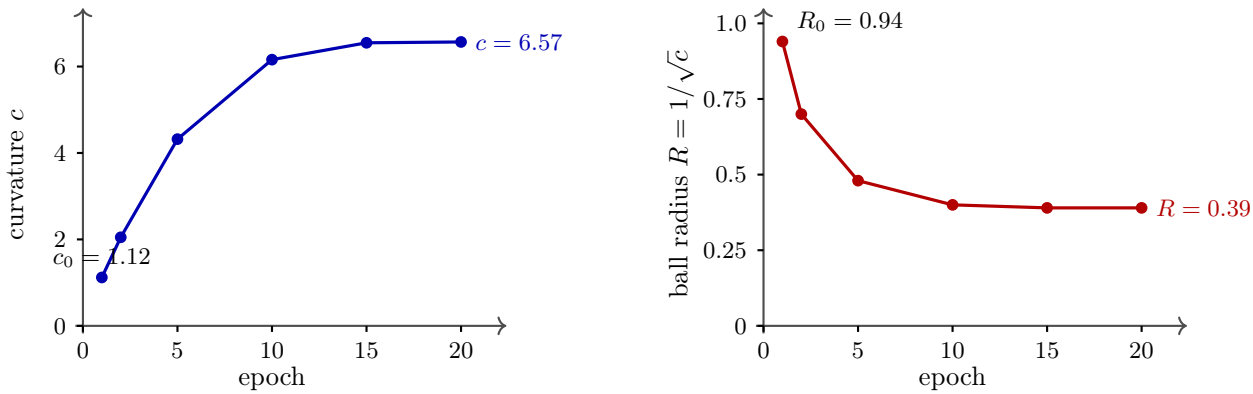

Figure 4: The learnable-curvature experiment on 10 M pairs. When the Poincaré-ball curvature parameter c is made trainable, gradient descent steadily increases it (left), which is equivalent to shrinking the ball's radius $R = 1/\sqrt{c}$ (right). Throughout the whole run, the embedding-norm standard deviation remains pinned at zero: every address sits exactly on the ball's boundary, regardless of where the boundary is. Given the option, the network prefers to shrink the ball toward a point rather than use its interior — which is the cleanest possible evidence that the hyperbolic geometry is the wrong prior for this task.

9 Discussion and Limitations

What the model is not. The reported metric approximates *static, daytime-average* OSRM times. It has no notion of live traffic, time of day, weather, road closures, or vehicle class. For production ETA it would have to be composed with a traffic-aware correction (which is the setting of most learned-ETA literature); for the interactive applications we target — exploratory real estate search, logistics clustering, reachability visualisation — the static metric is usually exactly what the user wants.

Error tails. The largest absolute errors consistently occur on a specific kind of trip: medium-distance routes across geographic barriers (estuaries in Brittany, the Rhône delta, the Alps). We read these as places where OSRM itself takes an unusual detour that our smooth metric cannot reproduce locally. A residual correction head conditioned on per-route features could plausibly halve the tail, at the cost of no longer being a pure metric.

Single tower vs. dual tower. We chose the symmetric single-tower model because, on the daytime-average labels used here, a dual-tower asymmetric variant offered no measurable gain. On labels that capture genuine asymmetry (live rush-hour data, one-way-heavy historical centres) dual towers would likely matter again. The loss of asymmetry is a deliberate trade-off, not a limitation of the approach.

The pivot-pruning cautionary tale. We emphasise Section 6's failure because it is easy to believe, in the abstract, that an index with a theoretically favourable intrinsic-dimension profile *must* outperform brute force. In practice, on geometrically tight embedding distributions, it does not: the savings from pruning are eaten alive by the per-candidate overhead, and GEMV on cached norms wins by being the simplest thing that saturates the memory bus. The lesson we draw is that for sub-100 ms queries over tens of millions of low-intrinsic-dim vectors, the right default is *brute force done well*, not a tree or pivot index that looks sophisticated on paper.

Scaling to other countries. The entire pipeline is reproducible per country: an OSM extract, an OSRM instance, a few days of spot-CPU time to generate pairs, and a retraining run. The model's hyperparameters (Space2Vec wavelengths, embedding dimension) are essentially country-agnostic.

10 Conclusion

We have shown that driving time over a country-scale road network can be accurately represented as Euclidean distance in a learned 64-dimensional space, and that the machinery required to make this useful — a Space2Vec front-end, a small residual backbone, a linear projection head that is in fact a learned Mahalanobis metric, a monotonic invertible mapping to seconds, and a brute-force vector database built around cached squared norms and the expansion identity $\|a-b\|^2 = \|a\|^2 + \|b\|^2 - 2\langle a, b\rangle$ — fits comfortably on commodity hardware. On ≈432 million OSRM pairs covering ∼33 M French addresses, the model reaches 1.40 min MAE and Spearman $\rho = 0.9989$, with 96.8% of predictions within five minutes of ground truth. The resulting embedding space is not only accurate: it has spontaneously reconstructed the topology of the French road network from travel-time supervision alone. Along the way, the failure of hyperbolic geometry on this task clarified what kind of metric space road networks actually want — a mesh, not a tree — and the failure of pivot-based pruning clarified how to build the corresponding vector index: by doing nothing clever, as fast as the memory bus allows.

Reproducibility. We intend to release a partial version of the training dataset, the model architecture and training script, and the RingDB index code under open licences. Full-corpus weights and the complete dataset are retained for commercial reasons.

References

[1] Hennie Daniels and Marina Velikova. Monotone and partially monotone neural networks. *IEEE Transactions on Neural Networks*, 21(6):906–917, 2010. doi:10.1109/TNN.2010.2044803. URL https://doi.org/10.1109/TNN.2010.2044803.

[2] Austin Derrow-Pinion, Jennifer She, David Wong, Oliver Lange, Todd Hester, Luis Perez, Marc Nunkesser, Seongjae Lee, Xueying Guo, Brett Wiltshire, Peter W. Battaglia, Vishal Gupta, Ang Li, Zhongwen Xu, Alvaro Sanchez-Gonzalez, Yujia Li, and Petar Veličković. ETA prediction with graph neural networks in Google Maps. In *Proceedings of the 30th ACM International Conference on Information & Knowledge Management (CIKM '21)*, pages 3767–3776, New York, NY, USA, 2021. ACM. doi:10.1145/3459637.3481916. URL https://doi.org/10.1145/3459637.3481916.

[3] Octavian-Eugen Ganea, Gary Bécigneul, and Thomas Hofmann. Hyperbolic neural networks. In *Advances in Neural Information Processing Systems 31 (NeurIPS 2018)*, pages 5350–5360. Curran Associates, Inc., 2018. URL https://arxiv.org/abs/1805.09112.

[4] Raia Hadsell, Sumit Chopra, and Yann LeCun. Dimensionality reduction by learning an invariant mapping. In *Proceedings of the 2006 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR '06)*, volume 2, pages 1735–1742. IEEE Computer Society, 2006. ISBN 0-7695-2597-0. doi:10.1109/CVPR.2006.100. URL https://doi.org/10.1109/CVPR.2006.100.

[5] Dennis Luxen and Christian Vetter. Real-time routing with OpenStreetMap data. In *Proceedings of the 19th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems (GIS '11)*, GIS '11, pages 513–516, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-1031-4. doi:10.1145/2093973.2094062. URL https://doi.org/10.1145/2093973.2094062.

[6] Gengchen Mai, Krzysztof Janowicz, Bo Yan, Rui Zhu, Ling Cai, and Ni Lao. Multi-scale representation learning for spatial feature distributions using grid cells. In *International Conference on Learning Representations (ICLR 2020)*, 2020. URL https://openreview.net/forum?id=rJljdh4KDH. Spotlight.

[7] Maximilian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations. In *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, pages 6338–6347. Curran Associates, Inc., 2017. URL https://arxiv.org/abs/1705.08039.

[8] Fatemeh Salehi Rizi, Jörg Schlötterer, and Michael Granitzer. Shortest path distance approximation using deep learning techniques. In *Proceedings of the 2018 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM '18)*, pages 1007–1014. IEEE, 2018. URL https://arxiv.org/abs/2002.05257.

[9] Florian Schroff, Dmitry Kalenichenko, and James Philbin. FaceNet: A unified embedding for face recognition and clustering. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 815–823. IEEE Computer Society, 2015. ISBN 978-1-4673-6964-0. doi:10.1109/CVPR.2015.7298682. URL https://doi.org/10.1109/CVPR.2015.7298682.

[10] Matthew Tancik, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, pages 7537–7547. Curran Associates, Inc., 2020. URL https://arxiv.org/abs/2006.10739.

[11] Valhalla Project Contributors. Valhalla: Open-source routing engine for OpenStreetMap. https://valhalla.github.io/valhalla/, 2026. Accessed: 2026-04-14.

[12] Mathias Vandaele. RingDB: A brute-force vector database for ring queries. https://github.com/mathias-vandaele/ring-db, 2026. Accessed: 2026-04-14.

[13] Antoine Wehenkel and Gilles Louppe. Unconstrained monotonic neural networks. In *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*, pages 1545–1555. Curran Associates, Inc., 2019. URL https://arxiv.org/abs/1908.05164.